

Qwen: base (few-shot) vs. LoRA (zero-shot, 1 epoch) (train dev_v2; test dev_v1; proper_term)

