




Qwen (base & LoRA): zero-shot vs. few-shot (train dev_v2; test dev_v1; proper_term)


Qwen 3B — Base (untrained)



Qwen 3B — LoRA (1 epoch)



Qwen 7B — Base (untrained)



Qwen 7B — LoRA (1 epoch)

