





Fine-Tuning Qwen2.5 on SAP dev_v2: LoRA Epoch Ablation vs GPT-4o-mini (train ~1,500 sent./lang · test held-out dev_v1 · proper_term)



Epoch 0: untuned Qwen base with 3-shot prompt; epochs 1-3: LoRA without few-shot. LoRA closes the BLEU gap on EN→ES/RU; GPT-4o-mini remains strongest for terminology accuracy.